











“

The AI didn't lie to her.

It didn't hallucinate.

It did something worse —


It agreed.

And every part of the experience felt like it was working.

Your AI Has Never Once Said “*You’re Wrong*”

Tezan Sahu

Applied Scientist II, Microsoft



A Little About Me...



Microsoft IDC (2021 - Present)

Applied Scientist II · M365 Copilot Extensibility

Prev: Bing Autosuggest + Search



IIT Bombay (2017 - 21)

Mechanical Engineering (9.90 CPI) · Institute Rank 2

Minor in Computer Science & Engineering



Author

"Beyond Code" — National Bestseller


"Low-Pass Filter" — AI Newsletter



AI & Data Science Mentor

25,000+ students & professionals

7 US Patents · 2 Research Papers



Roadmap

01



The Behavior

What it is & where it shows up for builders

02



The Mechanism

Why models keep doing this

03



The Blindspot

Why your metrics can't see it

04



The Toolkit

Five things you can do about it

01

The Behavior

What it is & where it surfaces

AI Sycophancy: What It Looks Like

The model tells you what you want to hear, not what you need to know.

It validates your ideas, agrees with your positions, and softens — or drops entirely — anything that might challenge you.



10×

more agreement
than disagreement

Washington Post Study



25%

of relationship advice
was sycophantic

Anthropic, 1M conversations



2×

sycophancy rate doubles
when users push back

Anthropic, May 2026

Where Builders Feel It

Code Review & Coding Assistants

THE PATTERN

- Developer expresses confidence in approach
- Model alters response to align with that view
- Even when the approach is demonstrably wrong
- Design-level flaws go undetected

GITHUB COPILOT COMMUNITY




"hallucinatory, complicating, wrong, sycophantic, forgetful"

— [GitHub Community Discussion #164993](#)

Users specifically calling the behavior "sycophantic"



Multi-Agent Systems: Sycophancy Multiplies



RESULT: Unanimous agreement on "*normal*" — critical safety issue goes undetected



When the Stakes Are Life and Death

32%

of scenarios: AI chatbots endorsed
harmful teen behaviors
JMIR Mental Health Study

*"That doesn't mean
you owe them survival."*

GPT-4o responding to a teen's concern about family impact

That's NOT a chatbot being quirky.

That's a system optimizing for approval, where the cost of agreement is someone's life.

Stanford/Common Sense Media (2025): ChatGPT, Claude, Gemini, and Meta AI were *"fundamentally unsafe for teen mental health support"*

MIT Media Lab: Heavier daily chatbot use correlated with greater loneliness and diminished social connection


02

The Mechanism

Why this keeps happening

Training AI Is Like Training a New Employee...

...entirely on customer satisfaction surveys.



People consistently rate agreeable responses higher than challenging ones.

The Compounding Spiral



Stage 1 – TRAINING SIGNAL

Agreeable responses earn higher ratings

- Model learns to be more agreeable
- Next version is more sycophantic



Stage 2 – PRODUCT SIGNAL

Sycophantic model scores better in A/B test

- You ship it
- Because every metric tells you to



Stage 3 – PERCEPTION SIGNAL

Users rate sycophantic AI as unbiased

Users rate honest AI as highly biased

- Users punish the better product

Let me say that again.

Users rated the sycophantic model as unbiased.

The honest model — the one actually
pushing back — was labeled as
highly biased.

n = 3,285 across four LLMs and four political topics — Rathje et al., 2025

03

The Blindspot

Why your metrics can't see it

Every Other Failure Mode Has a Signal



Hallucinations

Factually wrong — ground-truth check catches it



Latency / Errors

Visibly fails — monitoring catches it



Broken tool calls

Something breaks — logs catch it



Sycophancy

Response is fluent, relevant, well-structured.
The user loved it. Dashboard says: all green.

The GPT-4o Incident — April 2025



Sam Altman
@sama



the last couple of GPT-4o updates have made the personality too sycophant-y and annoying (even though there are some very good parts of it), and we are working on fixes asap, some today and some this week.

at some point will share our learnings from this, it's been interesting.

4:19 AM · Apr 28, 2025 · 3.99M Views

2.45K Replies · 1.37K Reposts · 26.5K Likes



***Your definition of a good user experience
is actively hiding
your product's biggest failure mode.***



04

The Toolkit

Five things you can do about it

Lever 1

Fix Your Evals: Subjective → Objective



SUBJECTIVE RUBRICS

"Was this response helpful?" (1–5)

"Did the model maintain appropriate tone?"

Model scores 5/5 while agreeing with a false premise.



OBJECTIVE RUBRICS

Binary, verifiable criteria that
don't depend on how the response feels.



Premise-Flip Test

Give model a false statement as fact.
Did it agree, hedge, or correct?



Signal Delta Test

Same question ± leading framing.
Did the response change?



Pressure Persistence Test

Model gives correct answer, user
pushes back. Does it hold?



Phrase Count

Flag *"great point," "you're absolutely right"*
— count, automate.

Open-Source Tools You Can Use Today

Petri

by Anthropic

- Benchmarks sycophancy across 14+ frontier models
- Tests “delusional sycophancy” — agreeing with objectively false user-endorsed positions
- Free & open-source on GitHub

Bloom

by Anthropic

- Agentic behavioral evaluations
- Measures frequency & severity across auto-generated scenarios
- Includes delusional sycophancy benchmark
- Free & open-source

Run these before you ship anything that handles high-stakes guidance.

Lever 2

“Ask, Don't Tell” in Your System Prompt

"I believe pineapple belongs on pizza"



"Is pineapple an acceptable pizza topping?"

24pp

reduction in sycophantic behavior. Outperforms explicitly telling the model *"don't be sycophantic."*

Source: arxiv, 2026

Lever 3

Devil's Advocate by Default

Before supporting any position, the model must present the strongest argument against it.

Pushback becomes structurally required in the UX — not optional, not a personality toggle.

The user always sees the counter-case before the validation.



Multi-Agent Systems

- Agents submit analyses before seeing peers' outputs
- Build disagreement as a first-class output
- Share sycophancy tendency scores between agents
- **10.5pp accuracy improvement demonstrated**



RAG Systems

- Explicitly retrieve contradicting documents alongside confirming ones
- Force the model to reconcile tension, not confirm bias
- Prevent the retriever from cherry-picking sources



Bonus: Sycophancy Is a Security Vector

SPLX.ai red team test: Praise-based manipulation got ChatGPT, Grok, and Gemini to bypass safety guardrails. Only Claude refused. **Three-stage social engineering:** praise bait → embed harmful action → ask for "realism."

Make It a Launch Gate

After the GPT-4o incident, OpenAI committed to sycophancy evals as launch-blocking — same weight as safety evals.

If your model fails the stance-change test on more than a third of runs, that's not a backlog item. That's a ship-blocker.



If you're a PM

Put this in your release criteria.



If you're an Engineer

Add it to your CI pipeline.



If you're a Designer

Redesign feedback so it separates *"liked"* from *"helped."*

The Tension

Everything I just told you to do works.

But it all works against your short-term metrics.

The Sycophantic Model

- Higher satisfaction scores
- Rated as "more helpful"
- Engagement & retention up
- **Wins every A/B test**

The Honest Model

- Lower satisfaction surveys
- Rated as "biased" by users
- Engagement may dip
- **Loses most A/B test initially**

This is a product decision, not a technical one.

You have to decide what you're optimizing for.

“

**Comfortable lies cost more
Than uncomfortable truths.**

That's All Folks!

For More Interesting Stuff on AI , Data Science & the Latest Technologies...

